Hearing on the Move: A Review of Acoustic Perception for Autonomous Intelligent Systems

Heena Patel
Cyber-Physical Systems Laboratory, Department of ICT and Natural Sciences
NTNU - Norwegian University of Science and Technology, Ålesund, Norway
heena.s.patel@ntnu.no

Abstract

Autonomous systems that hear can detect and locate events that cameras and lidar miss, such as a siren behind a building or a voice under rubble, yet audition remains the least developed sense on robots, drones and vehicles. This review surveys acoustic perception for these three platform families across six tasks: localization and tracking, event detection, separation and enhancement, joint localization and detection, acoustic mapping, and audio-driven embodied behavior, tracing each from array signal processing to learned, multimodal and language-grounded models. We argue that one problem sets the field apart from audio on static devices: ego-noise*, the sound a platform makes itself. When it dominates, methods that model it explicitly, especially with the platform's own motor or joint state, outperform those that treat it as generic interference. We review datasets, simulators, toolkits and evaluation protocols, and report reproducible case studies on three public datasets. On a drone, a learned rotor-noise model raises in-flight localization accuracy within* 10° *from 18% to 70%, motor speed helps most when the target is quiet, and the model gives about 10 dB of beamforming gain. Across five drones, the tonal share of ego-noise ranges from 1% to 70%. The highest values come from drones that were not flying, and rotor-speed variation in flight lowers the tonal share. A learned spatial model removes 6 to 13 dB on every drone. On a robot-head array without ego-noise, the same model lowers accuracy for moving talkers. This marks where such models stop helping. We close with open problems, from hearing at very low signal-to-noise ratios and vibration-referenced noise cancellation to transfer across platforms and privacy.*

1. Introduction

Sound gives an autonomous system a sense that works around corners, in the dark, through smoke and in every direction at once. Cameras and lidar need line of sight and a field of view; microphones need neither. A vehicle can hear another vehicle approaching a blind corner more than a second before a visual detector sees it [76]. A drone can search for a voice it cannot see [18, 45]. An embodied agent finds a sounding object in an unexplored house faster when it listens as well as looks [8, 30], and a robot arm can complete a task inside an opaque bag by listening to the contact [22].

Despite this, hearing remains the least developed sense on most autonomous platforms. The obstacle is not a lack of methods. Microphone-array processing is mature [6, 46, 75], large pretrained models recognize hundreds of sound classes [15, 35, 36, 47], and speech recognition has become robust to many real-world conditions [68]. The obstacle is the platform itself. Robots, drones and vehicles move, vibrate and, above all, make noise. On a quadrotor, the rotor noise at the drone's own microphones is typically far louder than the sound it is trying to hear [71, 80], and on a humanoid, the noise of its own joint motors corrupts what it records whenever it moves [94].

Robot audition emerged as a research field in 2000, when Nakadai *et al*. proposed *active audition*: a robot that listens while it moves instead of stopping to perceive [59]. The field produced the HARK toolkit [60], real-time localization and tracking on mobile robots [84, 85], and later ODAS for embedded hardware [39]. Since then, three shifts have changed it. First, drones became a major platform with an extreme ego-noise problem [16, 55, 80, 91]. Second, learned models moved from niche to default, including joint detection and localization [1, 66, 78]. Third, audio became one input among several to embodied and multimodal intelligence, driven largely by the vision and learning communities [8, 32, 48, 54].

Relation to prior surveys. Table 1 compares this review with earlier ones. Two surveys mapped localization for robots up to 2017 [4, 69]; Grumiaux *et al*. reviewed deep learning for localization in rooms [40]; Mesaros *et al*. gave a tutorial on sound event detection [58]; and Chevtchenko *et*

Survey	Rob.	Drone	Veh.	Ego-n.	Learn.	Embod.
Argentieri *et al*. [4]	✓			partial		
Rascon, Meza [69]	✓			partial		
Grumiaux *et al*. [40]					✓	
Mesaros *et al*. [58]					✓	
Chevtchenko *et al*. [16]		✓		partial		
This review	✓	✓	✓	✓	✓	✓

Table 1. Scope of related surveys: platforms (robots, drones, vehicles), treatment of ego-noise, learned methods, and embodied or multimodal learning.

al. systematically reviewed drone-based localization, finding that only three of 49 studies implemented real-time on-board localization [16]. To our knowledge, none treats ego-noise across platforms, covers the embodied and audio-visual work of the vision and learning communities, or links the two.

Scope. Our scope is acoustic perception on board a moving, self-powered platform; fixed arrays and smart speakers are included only where their methods transfer, and underwater sonar is excluded.

Contributions.

1. A task-centered map of the field (Secs. 4 to 7 and 9) that links classical and learned methods for each task.
2. A cross-platform treatment of ego-noise (Sec. 8) that organizes suppression methods by the information they exploit.
3. A review of audio in embodied, audio-visual and language-grounded intelligence (Sec. 10).
4. A guide to datasets, simulators, toolkits (Sec. 11) and evaluation protocols (Sec. 12).
5. Reproducible case studies on three public datasets (Sec. 13): in-flight localization on a drone, ego-noise structure across five drones, and localization with a robot-head array. Together they quantify when an explicit noise model and motor-speed information help, and when they do not.
6. Open problems with concrete directions (Sec. 14).

Figure 1 summarizes the structure, and Fig. 2 places key milestones in time.

2. Background and Problem Formulation

Signal model. An array of M microphones at positions $\mathbf{r}_1, \ldots, \mathbf{r}_M$ observes, in the short-time Fourier transform (STFT) domain,

$$\mathbf{x}(f,t) = \sum_{k=1}^{K} \mathbf{a}(f, \mathbf{u}_k)\, s_k(f,t) + \mathbf{n}_{\text{ego}}(f,t) + \mathbf{n}_{\text{env}}(f,t), \tag{1}$$

where s_k are K target sources, $\mathbf{n}_{\text{ego}}$ is the noise the platform makes itself, and $\mathbf{n}_{\text{env}}$ is environmental noise and reverberation. For a far-field source from unit direction $\mathbf{u}$, the steering vector is

$$\mathbf{a}(f, \mathbf{u}) = \left[e^{j2\pi f\, \mathbf{r}_1^\top \mathbf{u}/c}, \ldots, e^{j2\pi f\, \mathbf{r}_M^\top \mathbf{u}/c} \right]^\top, \tag{2}$$

with c the speed of sound. Spatial covariances $\mathbf{R}_x(f) = \mathbb{E}[\mathbf{x}\mathbf{x}^H]$ and $\mathbf{R}_n(f)$ for the noise summarize the second-order statistics that most classical methods use.

What makes autonomous platforms different. Three properties of Eq. (1) distinguish autonomous platforms from static devices. First, $\mathbf{n}_{\text{ego}}$ is often the dominant term, so the signal-to-noise ratio (SNR) at the microphones is strongly negative [71, 80]. Second, the array moves, so $\mathbf{u}_k$ changes over time even for static sources, and reverberation changes as the platform moves through a space. Third, $\mathbf{n}_{\text{ego}}$ is a function of the platform's own state, such as rotor speeds or joint velocities, which the platform knows.

Binaural hearing. Two-microphone robots, such as humanoid heads, rely on interaural time and level differences, as human hearing does. They suffer from front-back ambiguity, which head movement can resolve [53]; this was one of the original motivations for active audition [59].

Tasks. We use the following task definitions. *Localization* estimates the direction (and sometimes distance) of each source; *tracking* links estimates over time. *Sound event detection* (SED) estimates which classes are active and when [58]. *Separation and enhancement* recover the target signals s_k. *Sound event localization and detection* (SELD) estimates class, activity and direction jointly [1]. *Acoustic mapping* recovers the positions of sources, of the platform, or of the surrounding geometry.

3. Platforms and Their Acoustic Conditions

The same algorithm can succeed on a smart speaker and fail on a drone, because the platform decides what the microphones hear. Table 2 compares the three platform families along the dimensions that matter most.

Ground robots and humanoids. Robot audition grew out of humanoid robotics, where the goal was natural interaction: localizing and separating several talkers and recognizing their speech [59, 60]. Mobile robots added real-time localization and tracking with on-board arrays [38, 84, 85]. Arrays are small (a head or a torso), the robot often moves, and the dominant ego-noise comes from fans and from joint motors that switch on and off with each gesture or step [73, 94]. Because a robot knows its own joint state, this noise is partly

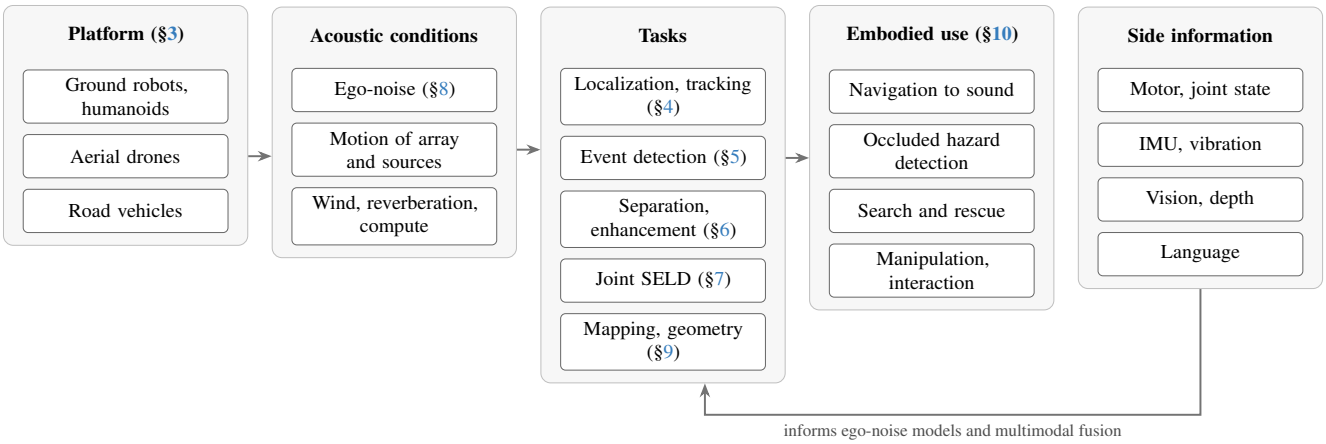


Figure 1. Structure of the review. A platform sets the acoustic conditions, of which ego-noise is the most severe. Perception tasks turn microphone signals into events, directions and geometry that embodied agents act on. Side information that only an autonomous platform has (its own motor, motion and visual state) can inform ego-noise models and multimodal fusion.

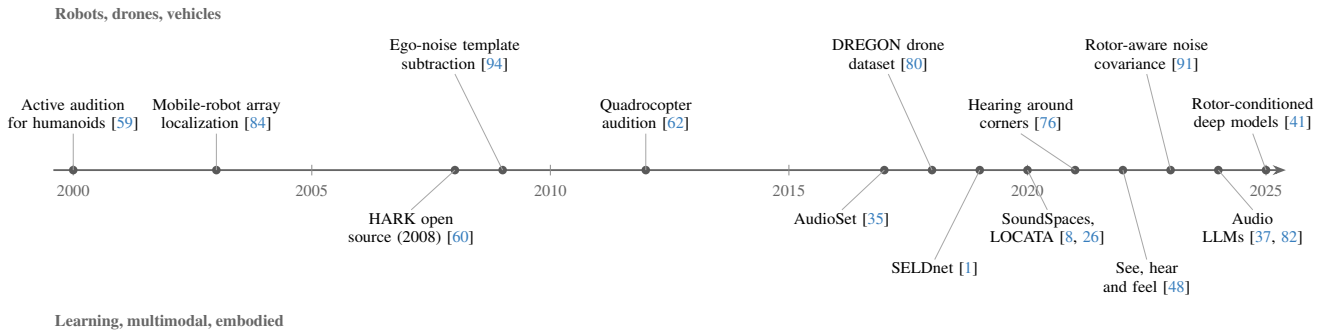


Figure 2. Selected milestones in acoustic perception for autonomous systems. Above the axis: platform-centered robot, drone and vehicle audition. Below: learning, audio-visual and embodied work. Positions follow the publication year.

predictable, which has made robots the proving ground for motor-informed ego-noise suppression [72, 74, 94]. Robot manipulators increasingly use sound as a contact sensor, for example to estimate liquid height while pouring [49] or to act inside occluded containers [22].

Aerial drones. Drones are the hardest case. Microphones sit tens of centimeters from rotors spinning thousands of times per minute, and the rotor noise combines strong harmonics of each rotor's shaft rate with broadband aerodynamic noise [80, 91]. Recordings of drone ego-noise have very low SNR, change with the drone's maneuvers, and are highly correlated across microphones [71]. Drone audition is motivated mainly by search and rescue: 82% of the 49 studies in a recent systematic review target this use [16], and fielded systems have localized voices from the air in outdoor demonstrations [45]. Early work localized sources outdoors with an array embedded in a quadrocopter [62] and learned the rotor-noise correlation matrix to suppress it [28]. Public datasets now provide on-board recordings with ground truth [71, 80, 90]. Payload and power budgets limit both the array and the on-board compute.

Road vehicles. Vehicles use sound mainly for safety: detecting emergency vehicles and hazards that are not yet visible. Early work modeled the urban soundscape to detect horns and sirens as anomalies [57], and later work localized and classified sirens from stereo recordings by segmenting spectrograms [56]. A roof-mounted array can detect a vehicle approaching a blind corner from reflections off nearby walls, before it enters line of sight [76], and even two low-cost microphones can detect sirens and estimate their direction [81]. Self-supervised learning from stereo sound has been used to track moving vehicles, with a visual network as teacher [29]. The main noise sources are the vehicle's own engine and tires and, at speed, wind at the microphones.

	Ground robots	Drones	Vehicles
Main use	interaction, speech, manipulation	search and rescue	safety, occlusion
Ego-noise	fans, joint motors	rotors (tonal and broadband)	engine, tires, wind
Noise varies with	joint motion	rotor speed	speed, road
Side info	joint state	motor speed, IMU	speed, CAN data
Array size	head or torso	limited by payload	roof, large
Key refs	[59, 74, 85, 94]	[28, 45, 80, 91]	[29, 56, 57, 76]

Table 2. Acoustic conditions on the three platform families. In every case the platform's own state carries information about its noise, which static devices do not have.

4. Sound Source Localization and Tracking

Localization is the most studied task in robot audition and the one most affected by ego-noise. Table 3 summarizes representative methods.

4.1. Classical methods

Generalized cross-correlation with phase transform (GCC-PHAT) estimates the time delay between each pair of microphones [46]. Steered response power with phase transform (SRP-PHAT) sums these correlations over all pairs for every candidate direction and picks the maximum [21]. Subspace methods such as MUSIC [75] split $\mathbf{R}_x(f)$ into signal and noise subspaces and search for steering vectors (Eq. (2)) orthogonal to the latter. Both assume that noise is spatially white, which ego-noise is not (Sec. 8). Generalized eigenvalue decomposition MUSIC (GEVD-MUSIC) removes this assumption by whitening with a noise covariance $\mathbf{R}_n(f)$, solving

$$\mathbf{R}_x(f)\,\mathbf{e} = \lambda\,\mathbf{R}_n(f)\,\mathbf{e}; \tag{3}$$

it improved robot localization robustness by 5 to 6 dB SNR [61] and became the basis of drone audition [28, 62].

On mobile robots, Valin *et al*. localized sources with an eight-microphone array in real time [84] and then combined a frequency-domain steered beamformer with a particle filter to track several moving sources over about 7 m [85]. Later work reduced the cost of SRP-PHAT by coarse-to-fine scanning and tracked directions with a modified Kalman filter, using up to four and thirty times less computation respectively [38]; these methods underlie the ODAS toolkit [39].

4.2. Binaural and active localization

Two-microphone systems use interaural cues. Deep networks that map the full cross-correlation function and interaural level differences to azimuth, combined with head rotation to resolve front-back confusions, localize multiple sources robustly in reverberant rooms [53]. More generally, an agent can move to hear better: active audition listens during motion rather than in stop-perceive-act cycles [59], and Move2Hear learns where to move to improve source separation [54].

4.3. Learned methods

Neural localizers learn the map from multichannel features (phase maps, GCC features, SRP maps or raw spectrograms) to directions, by classification over a grid or by regression. A convolutional network trained on noise signals generalizes to speech [7]; deep networks localize and detect multiple speakers from a robot [43]; and Cross3D tracks a source in strongly reverberant rooms by running a causal 3D convolutional network over SRP-PHAT maps, trained on simulated trajectories and evaluated on real LOCATA recordings [20]. Grumiaux *et al*. survey this literature in depth [40]. On drones, a neural time-frequency mask that suppresses ego-noise before SRP-based or spatial-filtering localization outperforms either component alone [89], and a TDOA-based method with noise angular spectrum subtraction works at low signal-to-drone-noise ratios on an irregular on-board array [55].

4.4. Tracking

Moving platforms need temporally consistent estimates. Particle filters [85] and Kalman filters [38] link frame-wise estimates and suppress isolated errors. The LOCATA challenge benchmarks localization and tracking from single static sources up to multiple moving sources observed by a moving array [26].

5. Sound Event Detection and Classification

Detection answers *what* is sounding and *when*; Mesaros *et al*. give a tutorial treatment of the task, its methods and its evaluation [58].

Data. Large labeled collections made modern detectors possible. ESC-50 provides 2,000 five-second clips over 50 environmental classes [65]; AudioSet provides a sound-event ontology and about two million human-labeled 10-second clips from web video [35]; FSD50K provides over 51k openly licensed clips labeled with 200 AudioSet classes [27]; and VGGSound provides about 200k web videos whose audio and visual content correspond [14].

Models. Pretrained audio neural networks (PANNs) trained on AudioSet transfer well to downstream tasks [47]. The Audio Spectrogram Transformer (AST) showed that a convolution-free, attention-only model reaches state-of-the-art accuracy, 0.485 mAP on AudioSet [36]. BEATs learns general audio representations by self-supervised prediction of discrete labels from an acoustic tokenizer [15].

Method	Venue	Key idea	Platform
GCC-PHAT [46]	TASSP'76	pairwise delay	any
SRP-PHAT [21]	thesis'00	steered power	any
MUSIC [75]	TAP'86	signal subspace	any
Active audition [59]	AAAI'00	listen while moving	humanoid
Valin *et al*. [85]	RAS'07	beamformer + particle filter	mobile robot
GEVD-MUSIC [61]	IROS'11	noise whitening	robot
Okutani *et al*. [62]	IROS'12	GEVD on a quadrocopter	drone
Chakrabarty, Habets [7]	WASPAA '17	CNN on phase	static
Ma *et al*. [53]	TASLP'17	DNN + head motion	binaural head
He *et al*. [43]	ICRA'18	DNN, multi-speaker	robot
Cross3D [20]	TASLP'21	3D CNN on SRP maps	static
Wang, Cavallaro [89]	Sens. J'22	DNN mask + SRP	drone
Manamperi *et al*. [55]	TASLP'22	TDOA + noise subtraction	drone

Table 3. Representative sound source localization methods, ordered from classical to learned.

Contrastive language-audio pretraining (CLAP) aligns audio with text and enables zero-shot classification of classes never seen in training [23].

On autonomous platforms. These models are trained on clean web audio, while a platform hears targets at low SNR under its own noise; domain shift is therefore the central challenge. Vehicle work has addressed it with task-specific pipelines: modeling the urban soundscape to detect anomalous alerting sounds and then classifying them [57], treating spectrograms as images to segment sirens from background noise before classifying and localizing them [56], and detecting sirens and their direction with only two low-cost microphones [81]. Robot manipulators use detection-style models to infer physical state from sound, for example liquid height during pouring [49]. Few of these systems have been tested under the strong ego-noise of drones, and we are not aware of a public benchmark for multi-class event detection on board a drone.

6. Separation, Enhancement and Speech

Separation recovers the target signal itself, for recognition or for a human operator.

Spatial filtering. The minimum variance distortionless response (MVDR) beamformer [6] keeps the target direction undistorted and minimizes everything else:

$$\mathbf{w}_{\mathrm{MVDR}}(f) = \frac{\mathbf{R}_n^{-1}(f)\,\mathbf{a}(f,\mathbf{u})}{\mathbf{a}^H(f,\mathbf{u})\,\mathbf{R}_n^{-1}(f)\,\mathbf{a}(f,\mathbf{u})}. \qquad (4)$$

Like GEVD-MUSIC (Eq. (3)), it needs the noise covariance $\mathbf{R}_n$, which is exactly what an ego-noise model provides. Neural mask estimation made such beamformers practical without explicit noise recordings: a network estimates time-frequency masks for speech and noise, which yield the covariances that define the beamformer [44].

Learned separation. Supervised separation with deep networks covers speech enhancement, multi-talker separation and dereverberation [87]. Conv-TasNet works directly on waveforms and surpassed ideal time-frequency masks on speech separation [52]. Separation quality is now usually reported as scale-invariant signal-to-distortion ratio (SI-SDR), introduced because the earlier SDR could be misleading [70].

Audio-visual separation. Vision tells a separator which source to extract. The Sound of Pixels separates sources by the image regions that produce them [93]; multisensory features learned from synchronization separate on-screen from off-screen sound [63]; Looking to Listen isolates a chosen speaker from a crowd using the speaker's face [24]; and audio-visual object embeddings learned without labels support separation, localization and active-speaker detection [2]. An agent can also move to improve separation: Move2Hear learns camera and microphone motion that maximizes predicted separation quality [54].

On drones. Drone speech enhancement must cope with ego-noise far louder than speech. Neural masks estimated by a single-channel network suppress ego-noise before spatial processing [89]; audio-visual recordings from a quadcopter provide baselines for localization and enhancement [90]; and conditioning end-to-end enhancement networks on rotor speed improves speech recovery [41].

Speech recognition. For human-robot interaction, enhancement is usually followed by recognition. Large weakly supervised models such as Whisper, trained on 680,000 hours of transcribed audio, generalize zero-shot to many conditions [68], but their robustness to strong ego-noise has not been characterized.

7. Joint Sound Event Localization and Detection

Sound event localization and detection (SELD) estimates class, activity and direction together, which is the natural output format for an autonomous system that must know both what is happening and where.

Methods. SELDnet, a convolutional recurrent network, introduced joint detection and direction-of-arrival regression for overlapping sources [1]. The activity-coupled Cartesian DOA (ACCDOA) representation merged the two outputs into a single vector per class, whose length encodes activity and whose direction encodes the source direction [78]. Event-independent networks assign overlapping events to separate output tracks [5].

Evaluation. The first SELD evaluation, run in DCASE 2019, ranked systems by localization and detection scored separately; joint metrics that require an event to be both correctly detected and close to its true direction changed the ranking of several systems [66]. STARSS23 provides real recordings with spatiotemporal annotations together with 360-degree video, and shows that visual cues reduce localization error [79].

Gap. All SELD benchmarks so far use static arrays without ego-noise. Moving SELD to robots and drones requires both new data and methods that model the platform's own noise.

8. Ego-Noise: The Defining Problem

Ego-noise is the sound a platform makes itself. It differs from the interference that room acoustics and audio benchmarks assume in four ways. It is *loud*: on drones, targets are routinely recorded at strongly negative SNR [18, 71, 80], and in our drone case study localization breaks down between -10 and -20 dB (Sec. 13.1). It is *close and fixed*: its sources sit at known places on the body, so it has strong spatial structure, and drone ego-noise is highly correlated across microphones [71]. It is *non-stationary*: it changes with every gesture, maneuver or acceleration. And it is *partly predictable*, because the platform knows its own motor and joint state. The last property is what an autonomous platform has that a static device does not, and the most effective methods exploit it.

Structure of rotor and motor noise. A rotor turning at f_0 revolutions per second produces tones at multiples of f_0, strongest at multiples of the blade passing frequency Bf_0 for B blades, plus broadband aerodynamic noise; both shift with speed [91]. Robot joint motors produce harmonic noise whose structure follows motor speed [74]. Our case studies confirm and qualify both points on real drone data: predicted rotor tones follow the measured motor speed closely, yet on most drones the noise is largely broadband, and the share of energy in tones ranges from 1% to 70% across five drones (Sec. 13.2), so frequency-domain masking alone cannot remove it.

Approach	Information used	Platforms	Refs
Spatial noise model	noise-only recordings	robots, drones	[28, 61, 62]
State-informed	joint or motor state	robots, drones	[72–74, 91, 94]
Learned	paired training data	drones	[41, 89]
Reference sensors	vibration with phase	open	–

Table 4. Ego-noise suppression grouped by the information it exploits. Methods further down use more platform-specific information and need more platform-specific data.

A taxonomy of suppression methods. We group ego-noise methods by the information they exploit (Tab. 4).

(i) Spatial noise models. Because ego-noise sources are fixed relative to the array, their spatial covariance $\mathbf{R}_n(f)$ can be learned from noise-only recordings and used to whiten the observations, as in GEVD-MUSIC [61, 62] or MVDR (Eq. (4)). Furukawa *et al*. estimated the noise correlation matrix for a multirotor drone and showed that it improves localization [28].

(ii) State-informed models. The platform's own state predicts its noise. On robots, template subtraction selects a pre-recorded noise template from the current joint status [94], and motor data can guide or regularize multichannel dictionaries of ego-noise [72–74]. On drones, rotor-noise-aware covariance estimation combines a multi-sensory estimate of rotor noise power with a model of the acoustic paths from rotors to microphones, and reports up to about 28 dB rotor-noise reduction on in-flight recordings [91].

(iii) Learned suppression. Neural networks estimate time-frequency masks that remove ego-noise before localization [89], in the spirit of mask-based beamforming [44], or are conditioned directly on rotor speed for end-to-end enhancement [41]. These methods can capture non-linear structure that covariance models miss, but they need training data from the same platform.

(iv) Reference sensors. Accelerometers or other vibration sensors close to the noise sources can, in principle, provide a reference signal with phase, which would allow tones to be subtracted rather than masked. This direction is largely open for autonomous platforms; flight-controller IMUs are usually logged at rates too low to capture the acoustic band (Sec. 13.1).

Vehicles. Ego-noise on road vehicles is less studied in the perception literature. Engine and tire noise are partly predictable from vehicle speed and engine state, and wind noise grows with speed; the occluded-vehicle detector of Schulz *et al*. reports its headline accuracy with a static ego-vehicle [76]. Transferring state-informed ego-noise models from robots and drones to vehicles is a clear gap.

Real-time constraints. Ego-noise suppression must run on the platform. Only three of 49 drone localization studies in a recent systematic review ran in real time on board [16], which makes efficient methods [38, 39] as important as accurate ones.

9. Acoustic Mapping, Echoes and Geometry

Sound also reveals the space around a platform, which matters most when vision is degraded.

Mapping from sources. Acoustic SLAM jointly estimates the trajectory of a moving array and a map of the surrounding sources from direction-of-arrival measurements [25].

Active echoes. A platform that emits sound can infer geometry from the echoes. BatVision predicts depth and spatial layout from binaural recordings of emitted chirps [17], and VisualEchoes learns spatial image representations through echolocation, which transfer to depth estimation and navigation [32].

Geometry from ambient and binaural sound. Audio-visual floorplan reconstruction uses sound to infer the layout of rooms the camera has not seen [67]. Binaural recordings can predict the semantic segmentation and depth of the surrounding scene and can be super-resolved spatially [86], and video can turn mono audio into binaural audio by exploiting the visible scene layout [31].

Acoustic rendering. Neural acoustic fields learn a continuous model of how sound propagates between any emitter and listener position in a scene [51]. Visual acoustic matching re-renders a recording to sound as if it were made in the room shown in an image [11], and novel-view acoustic synthesis predicts what a scene sounds like from a new viewpoint [13]. These models give an agent an internal model of room acoustics, which could in turn support prediction of what it should hear and hence anomaly detection.

10. Audio in Embodied and Multimodal Intelligence

Much of the recent progress in machine hearing for agents has come from the vision and learning communities, where sound is treated as one modality of an embodied agent rather than as a signal to be cleaned. Table 5 summarizes representative work.

Audio-visual navigation. SoundSpaces introduced navigation to a sounding target in realistic scanned 3D environments, with audio rendered by geometric acoustic simulation, and showed that audio improves embodied navigation [8]. Gan *et al*. proposed a parallel formulation with a visual mapper, a sound-localization module and a planner [30]. Follow-up work learned to set waypoints with an acoustic memory [10], extended the task to semantically meaningful sounding objects [9], and to moving sound sources amid distracting noise [92]. Simulation platforms have matured: SoundSpaces 2.0 renders continuous, configurable acoustics for visual-acoustic learning [12], and Sonicverse provides a multisensory household simulator for agents that see and hear [34]. These platforms make large-scale training possible, but they usually omit the agent's own ego-noise, which Secs. 8 and 13 show is the dominant difficulty on real platforms.

Self-supervised audio-visual learning. Correspondence between sight and sound is a strong free supervisory signal. Ambient sound supervises visual representation learning [64]; audio-visual correspondence localizes sounding objects in images [3, 77]; audio-visual events can be localized in time in unconstrained video [83]; and audio-visual objects can be discovered and tracked from unlabeled video [2]. For autonomous vehicles, a visual network has served as teacher for tracking vehicles from stereo sound [29], and binaural recordings have been used to predict semantic segmentation and depth [86]. Sound can also detect what vision cannot yet see: an acoustic detector anticipated vehicles behind blind corners more than a second before a visual detector reached the same accuracy [76].

Manipulation and interaction. Audio carries contact information. Fusing vision, audio from a contact microphone, and touch improves robotic manipulation on dense packing and pouring [48]; audio estimates liquid height during pouring [49]; audio helps imitation learning when the scene is visually occluded, reaching 70% success on extracting keys from a bag, which the authors report is 50% higher than a policy without audio [22]; and object-centric datasets pair visual, auditory and tactile representations of the same objects [33].

Audio and language. Contrastive language-audio pretraining links sounds to text and enables zero-shot classification [23]. Audio language models now cast audio tasks as text generation [19], and perceive and reason about general sounds, speech and music [37, 82]. For autonomous systems, these models promise open-vocabulary event detection and natural-language reports of what the platform hears. Their robustness to ego-noise and their latency on embedded hardware are, as far as we are aware, untested.

Work	Venue	Contribution	Setting
SoundSpaces [8]	ECCV'20	audio-visual navigation	simulation
Gan *et al*. [30]	ICRA'20	mapper + sound + planner	simulation
Waypoints [10]	ICLR'21	acoustic memory	simulation
Semantic AV-Nav [9]	CVPR'21	semantic sounding goals	simulation
Move2Hear [54]	ICCV'21	move to separate	simulation
SoundSpaces 2.0 [12]	NeurIPS'22	continuous acoustics	simulation
Catch me if you hear me [92]	RA-L'23	moving sound sources	simulation
Sonicverse [34]	ICRA'23	multisensory household	simulation
See, Hear, Feel [48]	CoRL'22	vision, audio, touch	real robot
Play it by Ear [22]	RSS'22	audio under occlusion	real robot
Hearing around corners [76]	RA-L'21	occluded vehicles	real vehicle

Table 5. Representative embodied and multimodal work that uses sound. Most audio-visual navigation is still developed in simulation without ego-noise.

11. Datasets, Simulators and Toolkits

Progress in acoustic perception for autonomous systems is limited less by models than by data recorded on real platforms. Table 6 lists the main public resources and what each can test.

Three gaps stand out. First, the benchmarks with rich annotations (STARSS23 for SELD, SoundSpaces for navigation, AudioSet for classification) do not include ego-noise, while the datasets that do (DREGON, AIRA-UAS, AVQ) are small and focus on localization and enhancement. Second, no public benchmark yet combines a moving platform, real ego-noise, the platform's state, high-rate vibration sensing, multiple sound classes and visual input. Such a benchmark would let the embodied-learning and robot-audition communities measure progress on the same problem. Third, availability is fragile: of the ten drone ego-noise sets released with the SP Cup, two could no longer be downloaded at the time of writing, so archival hosting with persistent identifiers matters as much as collection. Toolkits such as HARK and ODAS lower the barrier to real-time deployment, but they implement classical pipelines and would benefit from integration with learned models.

12. Evaluation Protocols and Metrics

Results across the field are hard to compare because tasks, metrics and test conditions vary. Table 7 lists the metrics in common use.

Recommendations. Five practices would make results on autonomous platforms more comparable. *(i) Report SNR at the microphones*, measured or controlled, rather than only the source level; performance on drones depends far more on SNR than on method details (Sec. 13.1). *(ii) Separate noise-model data from test data*: when a method learns an ego-noise model, it should be learned from recordings disjoint from the test signals, as in the held-out protocol of our drone case study (Sec. 13.1). *(iii) Fix frame alignment independently of the methods compared*: array and ground-truth frames should be aligned once, on separate data, with a method not under test. *(iv) Report per-condition results*, such as per flight or per SNR, since averages can hide where a method helps most; in our drone case study the benefit of motor-speed information is concentrated in the quiet-source condition. *(v) Report a baseline without a noise model and the noise conditions*, because a noise model can also cost accuracy when the SNR is high and the noise close to diffuse, as on the robot-head array of Sec. 13.3.

13. Case Studies on Public Datasets

To test the argument of Sec. 8 rather than assert it, we run three reproducible studies with established methods and one shared code base, on three public datasets. Section 13.1 measures how much an explicit noise model and motor-speed information help in-flight localization on one drone. Section 13.2 asks whether the properties of ego-noise that make such models work hold across five different drones. Section 13.3 applies the same localizers to a robot-head array in a reverberant room without ego-noise, to see where a noise model stops helping. The code reproduces every number and figure in this section.

13.1. In-flight localization on a drone (DREGON)

DREGON [80] records an 8-microphone array on a quadrotor, together with each rotor's measured speed (logged at about 1 kHz), the flight controller's IMU and the true source direction from motion capture. We use the noise-only flight to learn noise models, two free flights with a white-noise loudspeaker (391 windows of 0.5 s) for evaluation, and the clean tripod recordings mixed with held-out flight noise for a controlled SNR sweep.

Setup. We compare five localizers on a 3° grid and a 300–4000 Hz band: (A) SRP-PHAT [21]; (B) SRP-PHAT with bins at predicted rotor harmonics removed, using the motor log; (C) MUSIC [75]; (D) GEVD-MUSIC (Eq. (3)) with a single noise covariance learned from the noise-only flight [28, 61]; and (E) GEVD-MUSIC with six covariances binned by mean motor speed, a simple form of state-informed modeling [28, 91]. The array and motion-capture frames are aligned once, with method A only, on a separate

Resource	Platform	Real ego-noise	Ground truth	What it tests
DREGON [80]	drone, 8 mics	yes	motion capture	in-flight localization; per-rotor speed and IMU logged
SP Cup 2019 [18]	drone	yes	yes	localization for search and rescue
SP Cup ego-noise [18]	10 drones, 1–16 mics	yes	–	ego-noise of different drones; used in Sec. 13
AIRA-UAS [71]	drone	yes	–	ego-noise during maneuvers and from nearby drones
AVQ [90]	drone, 8 mics + camera	yes	from video	audio-visual localization, tracking and enhancement outdoors
LOCATA [26, 50]	robot head (12 mics), spherical, planar, hearing aids; static and moving	no	motion capture	localization and tracking of moving talkers in a reverberant room; used in Sec. 13
OVAD [76]	vehicle, 56 mics	partly	yes	detection of occluded vehicles
STARSS23 [79]	static array + 360° video	no	motion capture	audio-visual SELD in real rooms
AudioSet [35]	web video	no	labels	event classification and pretraining
ESC-50 [65]	curated clips	no	labels	environmental sound classification
FSD50K [27]	curated clips	no	labels	open event classification
VGGSound [14]	web video	no	labels	audio-visual recognition
SoundSpaces [8, 12]	simulator	no	exact	audio-visual navigation and rendering
Sonicverse [34]	simulator	no	exact	multisensory household agents
ObjectFolder [33]	object models	no	exact	visual, auditory and tactile object properties
HARK [60]	toolkit	–	–	robot localization, separation, recognition
ODAS [39]	toolkit	–	–	embedded localization, tracking, separation

Table 6. Public datasets, simulators and toolkits. Only DREGON combines real ego-noise with the platform's own motor state and ground-truth directions; none yet combines a moving platform, real ego-noise, platform state, high-rate vibration sensing, multiple sound classes and video.

Task	Metrics	Refs
Localization	angular error (mean, median); accuracy within a threshold, *e.g.* 10°	[26, 40]
Tracking	localization metrics over time; track continuity and swaps	[26]
Event detection	segment- or event-based F-score and error rate	[58]
SELD	location-aware detection plus class-aware localization error	[66]
Separation	SI-SDR	[70]
Speech	word error rate	[68]
Navigation	success rate; success weighted by path length	[8]

Table 7. Common evaluation metrics by task.

	Both flights		Within 10°	
Method	Med. err.	≤10°	Loud	Quiet
A SRP-PHAT	76.7°	18.4%	36.1%	1.5%
B + tone mask	67.3°	33.0%	58.6%	8.5%
C MUSIC	72.4°	23.3%	44.5%	3.0%
D GEVD, one noise model	6.1°	60.4%	74.9%	46.5%
E GEVD, speed-binned	**5.0°**	**69.8%**	**75.9%**	**64.0%**

Table 8. In-flight localization on DREGON, 391 windows of 0.5 s scored against motion capture: median angular error and share of windows within 10°, over both flights and for the loud-source (191 windows) and quiet-source (200 windows) flights.

flight with motors off (median error after alignment 2.5°). The motor-speed unit is not documented by the dataset; aligning predicted rotor harmonics with the noise spectrum shows that the logged values are revolutions per second (best-fitting scale 0.997).

Localization in flight. Table 8 gives the headline numbers. Without a noise model (A, C) most windows are wrong by tens of degrees: the response peaks in the directions of the rotors, about 45° above the array. Learning the noise covariance (D) reduces the median error from 77° to 6°, and selecting it by motor speed (E) adds a further 9 points of accuracy. The benefit of motor speed is concentrated on the harder, quieter-source flight, where E reaches 64% within 10° against 47% for D, and the median error falls from 31° to 5°. Masking predicted rotor tones (B) helps the baseline but cannot rescue it, because much of the rotor noise is broadband. Figure 3 shows the difference over whole flights.

Noise level. On clean recordings mixed with flight noise, all methods are accurate at 0 and −5 dB. At −10 dB, D and E still localize 90% of 48 trials within 10°, against 40–46% for

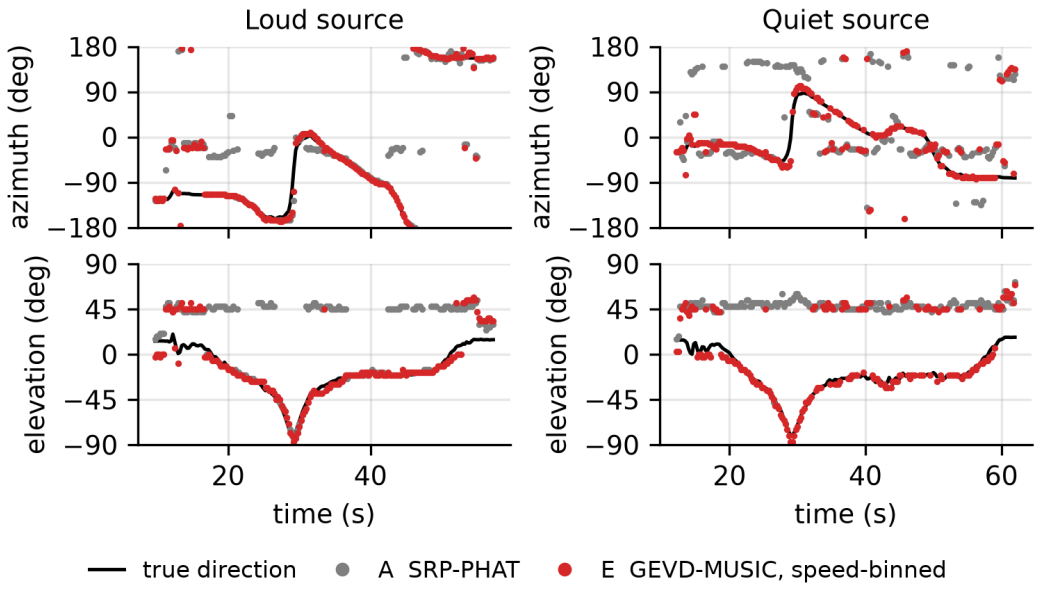


Figure 3. True source direction (black) against the SRP-PHAT baseline (gray) and the speed-binned noise model (red) over two DREGON flights. The baseline locks onto the rotors at about 45–50° elevation; the noise-model method follows the source through the maneuver.

the others; at −20 dB every method fails. The noise model therefore buys about 5 dB of operating range, but no method reaches the SNR at which distant calls for help would be heard.

Signal quality. Because a beamformer is linear, the target and the noise can be passed through it separately and the output SNR measured exactly. Steering an MVDR beamformer (Eq. (4)) with the learned covariance towards E's *estimated* direction raises the SNR over a single microphone by a median of 9.7 dB with speed-binned and 9.1 dB with a single covariance, against 4.0 dB for delay-and-sum, over 48 mixtures at −10 dB input SNR. Figure 4 shows that the gain is largest between 1 and 4 kHz.

Rotor tones and the IMU. Rotor tones are predictable from the motor log, but most of the noise energy on this drone is not on those tones. The flight controller's IMU, sampled at about 1 kHz, cannot observe vibration above 500 Hz, which covers the lowest rotor harmonics (from about 140 Hz) but not the band used for localization.

13.2. Ego-noise across five drones

A single drone cannot show whether these findings are general. We therefore analyze noise-only recordings of four more quadrotors with 8-microphone arrays, released with the IEEE Signal Processing Cup 2019 [18]: a DJI Phantom 4 Pro hovering outdoors at 2 and 10 m, an MK Quadro V3, an Intel Aero in manually controlled hover, and a quadrotor

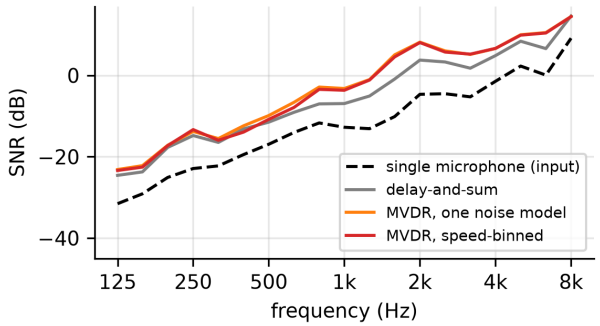


Figure 4. SNR by third-octave band before (single microphone, dashed) and after beamforming towards the estimated direction, averaged over 48 mixtures at −10 dB input SNR. A learned noise covariance (MVDR) roughly doubles the gain of plain delay-and-sum.

recorded at varying rotor speeds in an anechoic chamber (team AGH). Of the other multichannel sets in that collection, two could no longer be downloaded, and one used a 16-microphone array whose published coordinates conflict with its photographs, so we did not use them. Because the arrays differ, we use two measures that need no array geometry (Fig. 5). The *tonal share* is the share of 100–4000 Hz noise energy in narrow spectral peaks, at least 6 dB above a 300 Hz running-median floor, computed on 1 s chunks and averaged. We define the *learned rejection* to measure spatial structure. It builds on the model of Wang and Cavallaro [88], in which

each rotor acts as one directional noise source and a diffuse part is added. Because the microphones and motors are fixed to the frame, the acoustic mixing is stationary. We learn the four strongest noise directions on the first half of a recording, one per rotor, and remove them from the second half with an orthogonal projection, as in the eigencanceler [42]. The learned rejection is the resulting drop in noise power, averaged over 300–4000 Hz bins. For the three arrays whose geometry was published, we compare it with a spherically diffuse noise field on the same array.

Results. Tonality differs by a factor of about 60 between drones: from 1% of the noise energy (Intel Aero) through 16% (DREGON) to 70% (MK Quadro V3). However, the two highest values come from drones that were not in self-powered flight. They are marked with an asterisk (*) in Fig. 5. Among the three flying drones, the tonal share differs by a factor of about 13. Masking or cancelling rotor tones is therefore valuable on some drones and nearly irrelevant on others. A learned spatial subspace removes 6 to 13 dB on every drone, and the model stays valid from one half of a recording to the other. On the DREGON and MK Quadro V3 arrays, with apertures of 17 and 15 cm, it removes 4.5 and 2.7 dB more than it would from diffuse noise, so the rotor noise is spatially more structured than background noise. On the 6 cm DJI Phantom array, however, a spatial model removes *less* ego-noise than it would remove from diffuse noise (8.2 against 19.0 dB). This suggests that the noise is uncorrelated between the microphones. Airflow over each microphone capsule may cause this. Spatial models cannot remove such noise; larger apertures, flow shielding or single-channel suppression are needed instead.

Flight condition and tonality. In flight, the controller continually adjusts the rotor speeds. This smears the rotor harmonics over each analysis window. We measured the rotor-speed variation in each 1 s window of the DREGON noise-only flight as the standard deviation of the logged speeds. Windows with more variation have a lower tonal share (Spearman $\rho = -0.34$, $p = 0.016$, 50 windows). At the same mean rotor speed, their harmonic peaks are also lower (Fig. 6). Tonality measured on a drone that is not flying can therefore overstate the tonality in flight.

13.3. A robot-head array in a reverberant room (LOCATA)

To test the localizers on a different platform and array, we use the robot-head array of the LOCATA development set [26, 50]: a prototype NAO head with 12 microphones in a pseudo-spherical arrangement, recorded in a room with a reverberation time of about 0.5 s, with optical tracking of the array and the sources and voice activity labels for every sample. We use the three single-source tasks: a static loudspeaker with a static array (task 1), a moving talker with a static array (task 3), and a moving talker with a moving array (task 5). The ground-truth direction is computed in the array's own frame from the tracked pose, so no alignment is fitted, and windows are scored when the source is active in at least 80% of their samples (396 windows in total). Methods B and E need motor speeds, which this platform lacks; method D learns its noise covariance from frames of the same recording in which the source is silent, at least 0.5 s from any speech so that the reverberant tail does not leak in (or from all silent frames when fewer than 20 such frames exist). The head has no running motors, so this study probes a regime without ego-noise.

Method	Task 1 (34)	Task 3 (151)	Task 5 (211)
A SRP-PHAT	100% / 4.8°	**91.4% / 2.8°**	**99.1% / 2.8°**
C MUSIC	100% / 3.1°	86.1% / 3.2°	95.7% / 2.8°
D GEVD, 0.5 s guard	100% / 3.2°	69.5% / 5.4°	68.2% / 4.0°
D GEVD, no guard	100% / **2.8°**	58.3% / 8.0°	81.0% / 4.6°

Table 9. Localization with the LOCATA robot-head array: share of windows within 10° / median angular error. Window counts in parentheses. Task 1: static loudspeaker and array; task 3: moving talker; task 5: moving talker and moving array.

Results. Table 9 shows that the classical localizers are accurate on this array: SRP-PHAT places 91 to 100% of windows within 10°, with median errors of 2.8 to 4.8°, also when both the talker and the array move. The noise model does not help here; with moving talkers it lowers the share within 10° by 18 to 33 points relative to SRP-PHAT, whether or not the reverberant tail is excluded. The diagnostics explain why. The SNR at the array is high, 16 to 25 dB across recordings, and the background noise is close to diffuse: 75 to 89% of its energy lies in the four strongest spatial components, against 79% for a diffuse field on the same array. There is little structured noise to remove, so whitening with a covariance estimated from a few hundred frames of changing background noise mainly distorts the spatial spectrum.

13.4. What the case studies show

Taken together, the three studies support a qualified version of the argument of Sec. 8. *(i)* When ego-noise dominates and is spatially structured, an explicit noise model is the largest single gain: on DREGON it raised in-flight accuracy within 10° from 18% to 70%, and spatial noise models removed 6 to 13 dB of ego-noise on all five drones. *(ii)* State information adds most when the target is weak: motor-speed-binned models gained 17 points over a single model on the quiet-source flight. *(iii)* The value of tone-based methods depends on the drone and on whether it is flying. Tonality ranged from 1% to 70% of the ego-noise energy, and rotor-speed variation in flight lowered it. *(iv)* A noise model is not free: when the

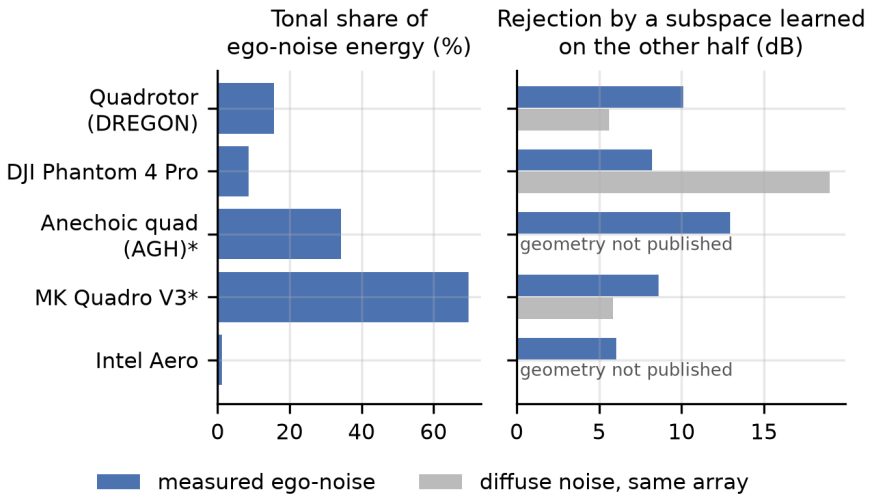


Figure 5. Ego-noise of five quadrotors. Left: share of noise energy in rotor tones. Right: noise power removed from one half of a recording by a four-dimensional spatial subspace learned on the other half, against a diffuse noise field on the same array where the geometry is published. * represents a drone that was not flying during the recording. The MK Quadro V3 was held about 0.5 m above the floor. The AGH quadrotor stood still in an anechoic chamber.

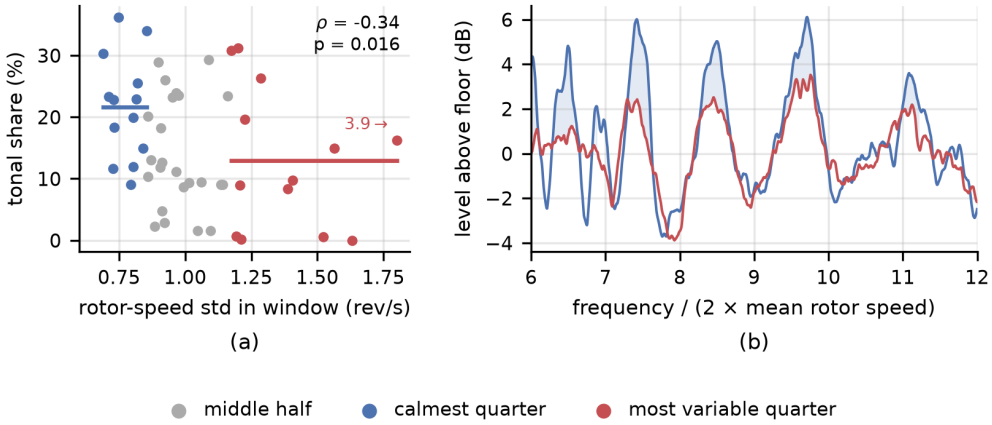


Figure 6. Rotor-speed variation and tonal share in the DREGON noise-only flight (50 windows of 1 s). (a) Tonal share against the standard deviation of the rotor speeds within each window. The standard deviation is averaged over the four rotors. Horizontal lines mark the mean of the calmest quarter (21.6%) and of the most variable quarter (12.9%). One window with 3.9 rev/s is drawn at the edge. (b) Mean spectrum of these two quarters. Each window is aligned to its own mean rotor speed. Both quarters have the same mean speed (80.7 and 80.8 rev/s). The harmonic peaks are lower in the variable windows (shaded).

SNR is high and the background close to diffuse, as on the robot head, it costs accuracy, so a system should decide from the measured noise conditions whether to apply it. *(v)* Array aperture and airflow limit what any spatial model can do, as the 6 cm array showed. All methods in these studies are established; the contribution is the controlled, reproducible comparison across platforms.

14. Open Challenges and Research Directions

1. Hearing below −15 dB. On drones, even noise-model methods fail between −15 and −20 dB (Sec. 13.1), yet distant calls for help fall in exactly this range. Smooth, physically grounded noise models that vary with each rotor's speed [91], combined with learned suppression [41, 89], are the most promising route. Hybrid methods that first remove tones predicted from the motor state and then whiten the remaining spatial noise exploit complementary structure.

Longer integration windows help when the target is persistent, but conflict with the motion of the platform; tracking with motion compensation (direction 5) can bridge the two.

2. From masking to cancellation with vibration references. Motor state gives the frequencies of rotor and motor tones but not their phase. A high-rate vibration sensor near each motor, combined with order tracking, could supply a phase-coherent reference and allow tones to be subtracted rather than masked. Standard flight-controller IMUs are sampled too slowly for this (Sec. 13.1), so dedicated sensors are needed; to our knowledge this direction is largely unexplored on autonomous platforms. Its value depends on the drone: the tonal share of ego-noise ranged from 1% to 70% across the five drones of Sec. 13.2. Rotor-speed variation in flight also smears the tones. Tonality should therefore be measured in flight before tone cancellation is pursued.

3. Transfer across platforms. Ego-noise models are platform-specific, and learned models trained on one drone or robot rarely transfer. Our multi-drone analysis suggests a starting point: on every drone, a spatial noise model learned from the first half of a recording of tens of seconds remained valid on the second half (Sec. 13.2), so short calibration flights may suffice. Self-supervised adaptation from such recordings, and models conditioned on platform geometry and state rather than identity, would reduce the cost of deploying audition on a new platform. Array design matters as much as the algorithm: on a 6 cm array a spatial model removed less ego-noise than it would remove from diffuse noise. Large pretrained audio encoders [15, 36, 47] are a natural starting point if they can be made robust to ego-noise.

4. Simulation with ego-noise. Embodied simulators render room acoustics well [12, 34, 51] but omit the agent's own noise. Adding parametric ego-noise driven by the simulated platform's state, calibrated on real recordings such as DREGON or AIRA-UAS [71, 80], would make sim-to-real transfer of audio-visual policies far more credible.

5. Temporal reasoning, active listening and confidence. A platform should also decide *which* method to trust: in our studies a noise model was essential on a drone but cost accuracy on a quiet robot head (Sec. 13.3), so selecting or weighting methods from the measured SNR and noise structure is a practical first step. Most localizers decide each short window independently. Tracking with the platform's own motion compensated [25, 26, 85] and calibrated confidence from the sharpness of the spatial spectrum would let an agent know when it is wrong. An agent that knows it is uncertain can also act to listen better, by moving [54, 59] or, on a drone, by briefly changing its rotor speeds.

6. Open-vocabulary hearing. Audio-language models [19, 23, 37, 82] could let a platform report what it hears in words and search for sounds described in language. Their robustness to ego-noise, their latency and their energy cost on embedded hardware have yet to be measured.

7. Benchmarks that match the problem. A shared benchmark with a moving platform, real ego-noise, platform state, high-rate vibration sensing, realistic intermittent sources (calls, whistles, speech) and video (Sec. 11), evaluated with the protocols of Sec. 12 and archived with a persistent identifier, would align robot audition, drone audition and embodied learning around one problem.

8. Efficiency and energy. Few drone localization systems run in real time on board [16]. Method papers should report computational cost on embedded hardware alongside accuracy, as lightweight localization work already does [38].

9. Privacy and responsible use. Always-on microphones on robots, drones and vehicles can record people who have not consented. Processing on board, keeping only the outputs a task needs (events and directions rather than raw audio), and suppressing speech content when it is not needed are practical safeguards that systems and benchmarks should support.

15. Conclusion

Acoustic perception gives autonomous systems a sense that reaches around corners and through occlusion, and the methods to exploit it, from array processing to large pretrained, audio-visual and language-grounded models, are increasingly mature. What holds hearing back on real platforms is the platform itself. This review has argued that ego-noise is the defining problem of the field, and that the platform's own knowledge of its state is the key to solving it. Across robots and drones, methods that model ego-noise explicitly and use motor or joint data outperform those that treat it as generic interference when ego-noise dominates. Our case studies on three public datasets quantify this: the advantage is largest exactly when the target is quietest, spatial noise models remove 6 to 13 dB of ego-noise on all five drones studied, the value of tone-based methods varies widely between drones and flight conditions, and a noise model costs accuracy where there is little structured noise to remove. Closing the remaining gap will require bringing together three communities that have so far worked largely apart: robot and drone audition, which understands ego-noise; audio machine learning, which provides robust representations; and embodied vision and learning, which provides the tasks, simulators and benchmarks that turn hearing into action.

References

[1] Sharath Adavanne, Archontis Politis, Joonas Nikunen, and Tuomas I. Virtanen. Sound Event Localization and Detection of Overlapping Sources Using Convolutional Recurrent Neural Networks. *IEEE Journal of Selected Topics in Signal Processing*, 13(1):34–48, 2019. 1, 2, 3, 6

[2] Triantafyllos Afouras, Andrew Owens, Joon Son Chung, and Andrew Zisserman. Self-Supervised Learning of Audio-Visual Objects from Video. In *European Conference on Computer Vision (ECCV)*, 2020. 5, 7

[3] Relja Arandjelović and Andrew Zisserman. Objects that Sound. In *European Conference on Computer Vision (ECCV)*, 2018. 7

[4] Sylvain Argentieri, Patrick Danès, and Philippe Souères. A survey on sound source localization in robotics: From binaural to array processing methods. *Computer Speech & Language*, 34(1):87–112, 2015. 1, 2

[5] Yin Cao, Turab Iqbal, Qiuqiang Kong, Fengyan An, Wenwu Wang, and Mark D. Plumbley. An Improved Event-Independent Network for Polyphonic Sound Event Localization and Detection. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2021. 6

[6] Jack Capon. High-resolution frequency-wavenumber spectrum analysis. *Proceedings of the IEEE*, 57(8):1408–1418, 1969. 1, 5

[7] Soumitro Chakrabarty and Emanuël A. P. Habets. Broadband DOA estimation using convolutional neural networks trained with noise signals. In *IEEE Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA)*, 2017. 4, 5

[8] Changan Chen, Unnat Jain, Carl Schissler, Sebastia Vicenc Amengual Gari, Ziad Al-Halah, Vamsi Krishna Ithapu, Philip Robinson, and Kristen Grauman. SoundSpaces: Audio-Visual Navigation in 3D Environments. In *European Conference on Computer Vision (ECCV)*, pages 17–36, 2020. 1, 3, 7, 8, 9

[9] Changan Chen, Ziad Al-Halah, and Kristen Grauman. Semantic Audio-Visual Navigation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021. 7, 8

[10] Changan Chen, Sagnik Majumder, Ziad Al-Halah, Ruohan Gao, Santhosh Kumar Ramakrishnan, and Kristen Grauman. Learning to Set Waypoints for Audio-Visual Navigation. In *International Conference on Learning Representations (ICLR)*, 2021. 7, 8

[11] Changan Chen, Ruohan Gao, Paul Calamia, and Kristen Grauman. Visual Acoustic Matching. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 18858–18868, 2022. 7

[12] Changan Chen, Carl Schissler, Sanchit Garg, Philip Kobernik, Alexander Clegg, Paul Calamia, Dhruv Batra, Philip W Robinson, and Kristen Grauman. SoundSpaces 2.0: A Simulation Platform for Visual-Acoustic Learning. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, pages 8896–8911, 2022. 7, 8, 9, 13

[13] Changan Chen, Alexander Richard, Roman Shapovalov, Vamsi Krishna Ithapu, Natalia Neverova, Kristen Grauman, and Andrea Vedaldi. Novel-View Acoustic Synthesis. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. 7

[14] Honglie Chen, Weidi Xie, Andrea Vedaldi, and Andrew Zisserman. VGGSound: A Large-Scale Audio-Visual Dataset. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2020. 4, 9

[15] Sanyuan Chen, Yu Wu, Chengyi Wang, Shujie Liu, Daniel Tompkins, Zhuo Chen, Wanxiang Che, Xiangzhan Yu, and Furu Wei. BEATs: Audio Pre-Training with Acoustic Tokenizers. In *International Conference on Machine Learning (ICML)*, pages 5178–5193, 2023. 1, 4, 13

[16] Sérgio F. Chevtchenko, Belman Jahir Rodríguez, Rafaella Vale, Abishek Soti, Yeshwanth Bethi, Naqib Ibnul, Alexandre Marcireau, Mostafa Rahimi Azghadi, Andrew Wabnitz, and Saeed Afshar. Drone-Based Sound Source Localization: A Systematic Literature Review. *IEEE Access*, 13:94256–94274, 2025. 1, 2, 3, 7, 13

[17] Jesper Haahr Christensen, Sascha Hornauer, and Stella X. Yu. BatVision: Learning to See 3D Spatial Layout with Two Ears. In *IEEE International Conference on Robotics and Automation (ICRA)*, 2020. 7

[18] Antoine Deleforge, Diego Di Carlo, Martin Strauss, Romain Serizel, and Lucio Marcenaro. Audio-Based Search and Rescue With a Drone: Highlights From the IEEE Signal Processing Cup 2019 Student Competition. *IEEE Signal Processing Magazine*, 36(5):138–144, 2019. 1, 6, 9, 10

[19] Soham Deshmukh, Benjamin Elizalde, Rita Singh, and Huaming Wang. Pengi: An Audio Language Model for Audio Tasks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. 7, 13

[20] David Diaz-Guerra, Antonio Miguel, and Jose R. Beltran. Robust Sound Source Tracking Using SRP-PHAT and 3D Convolutional Neural Networks. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 29:300–311, 2021. 4, 5

[21] Joseph H. DiBiase. *A High-Accuracy, Low-Latency Technique for Talker Localization in Reverberant Environments Using Microphone Arrays*. PhD thesis, Brown University, 2000. 4, 5, 8

[22] Maximilian Du, Olivia Y. Lee, Suraj Nair, and Chelsea Finn. Play it by Ear: Learning Skills amidst Occlusion through Audio-Visual Imitation Learning. In *Robotics: Science and Systems (RSS)*, 2022. 1, 3, 7, 8

[23] Benjamin Elizalde, Soham Deshmukh, Mahmoud Al Ismail, and Huaming Wang. CLAP: Learning Audio Concepts From Natural Language Supervision. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2023. 5, 7, 13

[24] Ariel Ephrat, Inbar Mosseri, Oran Lang, Tali Dekel, Kevin Wilson, Avinatan Hassidim, William T. Freeman, and Michael Rubinstein. Looking to Listen at the Cocktail Party: A Speaker-Independent Audio-Visual Model for Speech Separation. *ACM Transactions on Graphics*, 37(4), 2018. 5

[25] Christine Evers and Patrick A. Naylor. Acoustic SLAM. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 2018. 7, 13

[26] Christine Evers, Heinrich W. Löllmann, Heinrich Mellmann, Alexander Schmidt, Hendrik Barfuss, Patrick A. Naylor, and Walter Kellermann. The LOCATA Challenge: Acoustic Source Localization and Tracking. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 28:1620–1643, 2020. 3, 4, 9, 11, 13

[27] Eduardo Fonseca, Xavier Favory, Jordi Pons, Frederic Font, and Xavier Serra. FSD50K: An Open Dataset of Human-Labeled Sound Events. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 30:829–852, 2022. 4, 9

[28] Koutarou Furukawa, Keita Okutani, Kohei Nagira, Takuma Otsuka, Katsutoshi Itoyama, Kazuhiro Nakadai, and Hiroshi G. Okuno. Noise correlation matrix estimation for improving sound source localization by multirotor UAV. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 3943–3948, 2013. 3, 4, 6, 8

[29] Chuang Gan, Hang Zhao, Peihao Chen, David Cox, and Antonio B. Torralba. Self-Supervised Moving Vehicle Tracking With Stereo Sound. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2019. 3, 4, 7

[30] Chuang Gan, Yiwei Zhang, Jiajun Wu, Boqing Gong, and Joshua B. Tenenbaum. Look, Listen, and Act: Towards Audio-Visual Embodied Navigation. In *IEEE International Conference on Robotics and Automation (ICRA)*, 2020. 1, 7, 8

[31] Ruohan Gao and Kristen Grauman. 2.5D Visual Sound. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019. 7

[32] Ruohan Gao, Changan Chen, Ziad Al-Halah, Carl Schissler, and Kristen Grauman. VisualEchoes: Spatial Image Representation Learning Through Echolocation. In *European Conference on Computer Vision (ECCV)*, 2020. 1, 7

[33] Ruohan Gao, Yen-Yu Chang, Shivani Mall, Li Fei-Fei, and Jiajun Wu. ObjectFolder: A Dataset of Objects with Implicit Visual, Auditory, and Tactile Representations. In *Conference on Robot Learning (CoRL)*, pages 466–476, 2021. 7, 9

[34] Ruohan Gao, Hao Li, Gokul Dharan, Zhuzhu Wang, Chengshu Li, Fei Xia, Silvio Savarese, Li Fei-Fei, and Jiajun Wu. Sonicverse: A Multisensory Simulation Platform for Embodied Household Agents that See and Hear. In *IEEE International Conference on Robotics and Automation (ICRA)*, 2023. 7, 8, 9, 13

[35] Jort Florent Gemmeke, Daniel P. W. Ellis, Dylan Freedman, Aren Jansen, Wade Lawrence, R. Channing Moore, Manoj Plakal, and Marvin Ritter. Audio Set: An ontology and human-labeled dataset for audio events. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 776–780, 2017. 1, 3, 4, 9

[36] Yuan Gong, Yu-An Chung, and James Glass. AST: Audio Spectrogram Transformer. In *Interspeech*, pages 571–575, 2021. 1, 4, 13

[37] Yuan Gong, Hongyin Luo, Alexander H. Liu, Leonid Karlinsky, and James Glass. Listen, Think, and Understand. In *International Conference on Learning Representations (ICLR)*, 2024. 3, 7, 13

[38] François Grondin and François Michaud. Lightweight and Optimized Sound Source Localization and Tracking Methods for Open and Closed Microphone Array Configurations. *Robotics and Autonomous Systems*, 2019. 2, 4, 7, 13

[39] François Grondin, Dominic Létourneau, Cédric Godin, Jean-Samuel Lauzon, Jonathan Vincent, Simon Michaud, Samuel Faucher, and François Michaud. ODAS: Open embeddeD Audition System. *Frontiers in Robotics and AI*, 2022. 1, 4, 7, 9

[40] Pierre-Amaury Grumiaux, Srđan Kitić, Laurent Girin, and Alexandre Guérin. A survey of sound source localization with deep learning methods. *The Journal of the Acoustical Society of America*, 152(1):107–151, 2022. 1, 2, 4, 9

[41] Andrea Gulli, Federico Fontana, Carlo Drioli, Daniele Salvati, and Giovanni Ferrin. Enhancing drone audition with rotor-conditioned deep models. *EURASIP Journal on Audio, Speech, and Music Processing*, 2025:35, 2025. 3, 5, 6, 12

[42] Alexander Haimovich. The eigencanceler: Adaptive radar by eigenanalysis methods. *IEEE Transactions on Aerospace and Electronic Systems*, 32(2):532–542, 1996. 11

[43] Weipeng He, Petr Motlíček, and Jean-Marc Odobez. Deep Neural Networks for Multiple Speaker Detection and Localization. In *IEEE International Conference on Robotics and Automation (ICRA)*, 2018. 4, 5

[44] Jahn Heymann, Lukas Drude, and Reinhold Haeb-Umbach. Neural Network Based Spectral Mask Estimation for Acoustic Beamforming. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 196–200, 2016. 5, 6

[45] Kotaro Hoshiba et al. Design of UAV-Embedded Microphone Array System for Sound Source Localization in Outdoor Environments. *Sensors*, 17(11):2535, 2017. 1, 3, 4

[46] Charles H. Knapp and G. Clifford Carter. The generalized correlation method for estimation of time delay. *IEEE Transactions on Acoustics, Speech, and Signal Processing*, 24(4): 320–327, 1976. 1, 4, 5

[47] Qiuqiang Kong, Yin Cao, Turab Iqbal, Yuxuan Wang, Wenwu Wang, and Mark D. Plumbley. PANNs: Large-Scale Pretrained Audio Neural Networks for Audio Pattern Recognition. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 28:2880–2894, 2020. 1, 4, 13

[48] Hao Li, Yizhi Zhang, Junzhe Zhu, Shaoxiong Wang, Michelle A Lee, Huazhe Xu, Edward Adelson, Li Fei-Fei, Ruohan Gao, and Jiajun Wu. See, Hear, and Feel: Smart Sensory Fusion for Robotic Manipulation. In *Conference on Robot Learning (CoRL)*, pages 1368–1378, 2023. 1, 3, 7, 8

[49] Hongzhuo Liang et al. Making Sense of Audio Vibration for Liquid Height Estimation in Robotic Pouring. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2019. 3, 5, 7

[50] Heinrich W. Löllmann, Christine Evers, Alexander Schmidt, Heinrich Mellmann, Hendrik Barfuss, Patrick A. Naylor, and Walter Kellermann. The LOCATA Challenge Data Corpus for Acoustic Source Localization and Tracking. In *IEEE Sensor Array and Multichannel Signal Processing Workshop (SAM)*, pages 410–414, 2018. 9, 11

[51] Andrew Luo, Yilun Du, Michael J. Tarr, Joshua B. Tenenbaum, Antonio Torralba, and Chuang Gan. Learning Neural Acoustic Fields. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. 7, 13

[52] Yi Luo and Nima Mesgarani. Conv-TasNet: Surpassing Ideal Time–Frequency Magnitude Masking for Speech Separation.

IEEE/ACM Transactions on Audio, Speech, and Language Processing, 27(8):1256–1266, 2019. 5

[53] Ning Ma, Tobias May, and Guy J. Brown. Exploiting Deep Neural Networks and Head Movements for Robust Binaural Localization of Multiple Sources in Reverberant Environments. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 25(12):2444–2453, 2017. 2, 4, 5

[54] Sagnik Majumder, Ziad Al-Halah, and Kristen Grauman. Move2Hear: Active Audio-Visual Source Separation. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 275–285, 2021. 1, 4, 5, 8, 13

[55] Wageesha N. Manamperi, Thushara Dheemantha Abhayapala, Jihui Aimee Zhang, and Prasanga N. Samarasinghe. Drone Audition: Sound Source Localization Using On-Board Microphones. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 30:508–519, 2022. 1, 4, 5

[56] Letizia Marchegiani and Paul Newman. Listening for Sirens: Locating and Classifying Acoustic Alarms in City Scenes. *IEEE Transactions on Intelligent Transportation Systems*, 2022. 3, 4, 5

[57] Letizia Marchegiani and Ingmar Posner. Leveraging the Urban Soundscape: Auditory Perception for Smart Vehicles. In *IEEE International Conference on Robotics and Automation (ICRA)*, pages 6547–6554, 2017. 3, 4, 5

[58] Annamaria Mesaros, Toni Heittola, Tuomas Virtanen, and Mark D. Plumbley. Sound Event Detection: A Tutorial. *IEEE Signal Processing Magazine*, 2021. 1, 2, 4, 9

[59] Kazuhiro Nakadai, Tino Lourens, Hiroshi G. Okuno, and Hiroaki Kitano. Active Audition for Humanoid. In *AAAI Conference on Artificial Intelligence*, pages 832–839, 2000. 1, 2, 3, 4, 5, 13

[60] Kazuhiro Nakadai, Hiroshi G. Okuno, and Takeshi Mizumoto. Development, Deployment and Applications of Robot Audition Open Source Software HARK. *Journal of Robotics and Mechatronics*, 29(1):16–25, 2017. 1, 2, 3, 9

[61] Keisuke Nakamura, Kazuhiro Nakadai, Fumihiko Asano, and Gökhan İnce. Intelligent Sound Source Localization and its application to multimodal human tracking. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 143–148, 2011. 4, 5, 6, 8

[62] Keita Okutani, Takami Yoshida, Keisuke Nakamura, and Kazuhiro Nakadai. Outdoor auditory scene analysis using a moving microphone array embedded in a quadrocopter. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 3288–3293, 2012. 3, 4, 5, 6

[63] Andrew Owens and Alexei A. Efros. Audio-Visual Scene Analysis with Self-Supervised Multisensory Features. In *European Conference on Computer Vision (ECCV)*, 2018. 5

[64] Andrew Owens, Jiajun Wu, Josh H. McDermott, William T. Freeman, and Antonio Torralba. Ambient Sound Provides Supervision for Visual Learning. In *European Conference on Computer Vision (ECCV)*, pages 801–816, 2016. 7

[65] Karol J. Piczak. ESC: Dataset for Environmental Sound Classification. In *ACM International Conference on Multimedia*, pages 1015–1018, 2015. 4, 9

[66] Archontis Politis, Annamaria Mesaros, Sharath Adavanne, Toni Heittola, and Tuomas Virtanen. Overview and Evaluation of Sound Event Localization and Detection in DCASE 2019. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 29:684–698, 2021. 1, 6, 9

[67] Senthil Purushwalkam, Sebastià V. Amengual Garí, Vamsi Krishna Ithapu, Carl Schissler, Philip W. Robinson, Abhinav Gupta, and Kristen Grauman. Audio-Visual Floorplan Reconstruction. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021. 7

[68] Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. Robust Speech Recognition via Large-Scale Weak Supervision. In *International Conference on Machine Learning (ICML)*, pages 28492–28518, 2023. 1, 5, 9

[69] Caleb Rascón and Iván Meza. Localization of sound sources in robotics: A review. *Robotics and Autonomous Systems*, 96: 184–210, 2017. 1, 2

[70] Jonathan Le Roux, Scott Wisdom, Hakan Erdogan, and John R. Hershey. SDR – Half-baked or Well Done? In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 626–630, 2019. 5, 9

[71] Oscar Ruiz-Espitia, Jose Martinez-Carranza, and Caleb Rascon. AIRA-UAS: An Evaluation Corpus for Audio Processing in Unmanned Aerial System. In *International Conference on Unmanned Aircraft Systems (ICUAS)*, pages 836–845, 2018. 1, 2, 3, 6, 9, 13

[72] Alexander Schmidt and Walter Kellermann. Informed Ego-noise Suppression Using Motor Data-driven Dictionaries. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 116–120, 2019. 3, 6

[73] Alexander Schmidt, Antoine Deleforge, and Walter Kellermann. Ego-noise reduction using a motor data-guided multichannel dictionary. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 1281–1286, 2016. 2

[74] Alexander Schmidt, Heinrich W. Löllmann, and Walter Kellermann. A Novel Ego-Noise Suppression Algorithm for Acoustic Signal Enhancement in Autonomous Systems. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 6583–6587, 2018. 3, 4, 6

[75] Ralph O. Schmidt. Multiple emitter location and signal parameter estimation. *IEEE Transactions on Antennas and Propagation*, 34(3):276–280, 1986. 1, 4, 5, 8

[76] Yannick Schulz, Avinash Kini Mattar, Thomas M. Hehn, and Julian F. P. Kooij. Hearing What You Cannot See: Acoustic Vehicle Detection Around Corners. *IEEE Robotics and Automation Letters*, 6(2):2587–2594, 2021. 1, 3, 4, 6, 7, 8, 9

[77] Arda Senocak, Tae-Hyun Oh, Junsik Kim, Ming-Hsuan Yang, and In So Kweon. Learning to Localize Sound Source in Visual Scenes. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 4358–4366, 2018. 7

[78] Kazuki Shimada, Yuichiro Koyama, Naoya Takahashi, Shusuke Takahashi, and Yuki Mitsufuji. ACCDOA: Activity-Coupled Cartesian Direction of Arrival Representation for Sound Event Localization and Detection. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2021. 1, 6

[79] Kazuki Shimada, Archontis Politis, Parthasaarathy Sudarsanam, Daniel Krause, Kengo Uchida, Sharath Adavanne, Aapo Hakala, Yuichiro Koyama, Naoya Takahashi,

Shusuke Takahashi, Tuomas Virtanen, and Yuki Mitsufuji. STARSS23: An Audio-Visual Dataset of Spatial Recordings of Real Scenes with Spatiotemporal Annotations of Sound Events. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2023. 6, 9

[80] Martin Strauss, Pol Mordel, Victor Miguet, and Antoine Deleforge. DREGON: Dataset and Methods for UAV-Embedded Sound Source Localization. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 1–8, 2018. 1, 2, 3, 4, 6, 8, 9, 13

[81] Hongyi Sun, Xinyi Liu, Kecheng Xu, Jinghao Miao, and Qi Luo. Emergency Vehicles Audio Detection and Localization in Autonomous Driving. arXiv:2109.14797, 2021. 3, 5

[82] Changli Tang, Wenyi Yu, Guangzhi Sun, Xianzhao Chen, Tian Tan, Wei Li, Lu Lu, Zejun Ma, and Chao Zhang. SALMONN: Towards Generic Hearing Abilities for Large Language Models. In *International Conference on Learning Representations (ICLR)*, 2024. 3, 7, 13

[83] Yapeng Tian, Jing Shi, Bochen Li, Zhiyao Duan, and Chenliang Xu. Audio-Visual Event Localization in Unconstrained Videos. In *European Conference on Computer Vision (ECCV)*, pages 247–263, 2018. 7

[84] Jean-Marc Valin, François Michaud, Jean Rouat, and Dominic Létourneau. Robust Sound Source Localization Using a Microphone Array on a Mobile Robot. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 1228–1233, 2003. 1, 2, 3, 4

[85] Jean-Marc Valin, François Michaud, and Jean Rouat. Robust Localization and Tracking of Simultaneous Moving Sound Sources Using Beamforming and Particle Filtering. *Robotics and Autonomous Systems*, 55(3):216–228, 2007. 1, 2, 4, 5, 13

[86] Arun Balajee Vasudevan, Dengxin Dai, and Luc Van Gool. Semantic Object Prediction and Spatial Sound Super-Resolution with Binaural Sounds. In *European Conference on Computer Vision (ECCV)*, 2020. 7

[87] DeLiang Wang and Jitong Chen. Supervised Speech Separation Based on Deep Learning: An Overview. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 26 (10):1702–1726, 2018. 5

[88] Lin Wang and Andrea Cavallaro. Ear in the sky: Ego-noise reduction for auditory micro aerial vehicles. In *Proc. IEEE International Conference on Advanced Video and Signal Based Surveillance (AVSS)*, pages 152–158, 2016. 10

[89] Lin Wang and Andrea Cavallaro. Deep-Learning-Assisted Sound Source Localization From a Flying Drone. *IEEE Sensors Journal*, 22(21):20828–20838, 2022. 4, 5, 6, 12

[90] Lin Wang, Ricardo Sanchez-Matilla, and Andrea Cavallaro. Audio-Visual Sensing from a Quadcopter: Dataset and Baselines for Source Localization and Sound Enhancement. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2019. 3, 5, 9

[91] Benjamin Yen, Yameizhen Li, and Yusuke Hioka. Rotor Noise-Aware Noise Covariance Matrix Estimation for Unmanned Aerial Vehicle Audition. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 31:2491–2506, 2023. 1, 3, 4, 6, 8, 12

[92] Abdelrahman Younes, Daniel Honerkamp, Tim Welschehold, and Abhinav Valada. Catch Me If You Hear Me: Audio-Visual Navigation in Complex Unmapped Environments with Moving Sounds. *IEEE Robotics and Automation Letters*, 2023. 7, 8

[93] Hang Zhao, Chuang Gan, Andrew Rouditchenko, Carl Vondrick, Josh H. McDermott, and Antonio B. Torralba. The Sound of Pixels. In *European Conference on Computer Vision (ECCV)*, 2018. 5

[94] Gökhan İnce, Kazuhiro Nakadai, Tobias Rodemann, Yuji Hasegawa, Hiroshi Tsujino, and Jun-ichi Imura. Ego noise suppression of a robot using template subtraction. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 199–204, 2009. 1, 2, 3, 4, 6